

# Safety-Gated, Age-Adaptive LoRA Fine-Tuning for a Virtual Pediatric Hospital Guide

Anthony Barros<sup>1\*</sup>, Kelsey Kloosterman<sup>1</sup>,  
Sandrine de Ribaupierre<sup>2,3</sup>, Roy Eagleson<sup>1,2</sup>

<sup>1\*</sup>Department of Artificial Intelligence and Software Engineering,  
Faculty of Engineering, Western University, 1151 Richmond Street,  
London, N6A3K7, Ontario, Canada.

<sup>2</sup>Centre for Brain and Mind, Western University, 1151 Richmond  
Street, London, N6A3K7, Ontario, Canada.

<sup>3</sup>Clinical Neurological Sciences, Western University, 1151 Richmond  
Street, London, N6A3K7, Ontario, Canada.

\*Corresponding author(s). E-mail(s): [abarros6@uwo.ca](mailto:abarros6@uwo.ca);  
Contributing authors: [kkloost2@uwo.ca](mailto:kkloost2@uwo.ca); [sderibau@uwo.ca](mailto:sderibau@uwo.ca);  
[eagleson@uwo.ca](mailto:eagleson@uwo.ca);

## Abstract

Pediatric patients need age-appropriate, institutionally grounded information that general-purpose AI systems are not designed to provide, and any system accepting free-text input from children must defend against inputs it was never designed to handle. We present *small-bear*, an age-targeted conversational guide deployed within a digital recreation of Victoria Hospital (London, Ontario), paired with our *guard-bear*, an upstream input-safety classifier – together a complete pipeline running entirely on consumer hardware (an Apple Mac Mini M4, 16 GB) with no cloud dependency. The first finding concerns safety: the unmodified base classifier we build on, **Prompt-Guard-86M** [1], performs *worse than random* (ROC-AUC 0.448) on pediatric-context queries, saturating by flagging nearly everything as unsafe. Fine-tuning *guard-bear* on 4,053 pediatric-specific labeled examples (deduplicated, verified disjoint train/val/test) resolves this to near-perfect discrimination, confirmed across a 5-seed retraining campaign (ROC-AUC **0.9992 ± 0.0005**, recall **0.988 ± 0.004**) – evidence that safety mechanisms validated for a developer-facing threat model do not transfer to a different population without dedicated evaluation. The second finding concerns *small-bear*: we fine-tune Llama 3.2 Instruct (1B, 3B) with LoRA at two configurations, Standard ( $r = 8$ , 16 layers) and Fast ( $r = 4$ , 8 layers), on 1123 LIMA-inspired [2] synthetic examples, substantially outperforming zero-shot base models on readability (56–66% FK  $\leq 7.0$  pass rate vs. 12%) and

inter-role lexical separation (0.92–0.96 vs. 0.64–0.67). A cheaper prompted-but-not-fine-tuned baseline matches or exceeds these metrics (92–98%, 0.970–1.000) but at 1.2–2× the tokens/response, failing the 1.0-second latency target it costs nothing to attempt; fine-tuning’s real contribution is decoupling register from length, with Fast-1B the only Llama configuration meeting that target on average for the 5–11 role. During development, an unexpected **configuration-ordering crossover** emerged (Standard better on 3B, Fast better on 1B); a dedicated 110-run multi-seed campaign confirmed it is statistically real on both sides, but not real in the same way: the 1B-side gap is a length-independent register effect, while the 3B-side gap, once length is controlled for, is only marginal ( $p = \mathbf{0.055}$ ) and substantially attributable to response length rather than register. Together, this pipeline’s components each run practically on consumer hardware (combined, end-to-end latency remains unmeasured), and its safety layer, not response quality, posed the larger practical risk when unvalidated.

**Keywords:** parameter-efficient fine-tuning, low-rank adaptation, pediatric clinical NLP, virtual reality, on-device inference, small language models, LLM input safety classification

## 1 Introduction

Large language models (LLMs) enable domain-specific conversational systems, but resource-constrained clinical deployment raises constraints rarely addressed together: limited memory favors smaller models, domain adaptation requires fine-tuning, real-time interaction imposes strict latency bounds, and free-text input from a vulnerable population must be defended against inputs the system was never designed to handle.

This paper presents *small-bear*, an age-targeted conversational guide for children aged 5–11 and adolescents aged 12–18 embedded in a VR recreation of Victoria Hospital (London, Ontario), together with *guard-bear*, an upstream safety classifier – both named after LHSC’s bear mascot. All training and inference runs on an Apple Mac Mini M4 (16 GB): consumer hardware, no cloud infrastructure.

A conversational system accepting open-ended text from children needs a defense layer distinct from *small-bear* itself. General-purpose prompt-injection classifiers such as **Prompt-Guard-86M** [1] target a developer-facing threat model that does not describe a pediatric system’s actual inputs: short, first-person legitimate queries on one side, gibberish and adult content on the other. Our *guard-bear* closes this gap, finding, as our central safety result, that the unmodified base classifier performs worse than random here – a failure fine-tuning fully resolves.

For *small-bear*, we fine-tune Llama 3.2 Instruct models [3] at 1B and 3B scales using LoRA [4], alongside QLoRA [5]. Training data follows the LIMA hypothesis [2]: ~560 examples per adapter (1123 total) from Victoria Hospital’s public materials. We compare Standard ( $r = 8, 16$  layers) and Fast ( $r = 4, 8$  layers) adapters across both sizes and age groups. Pediatric deployment demands developmental vocabulary, supportive register, and hard safety constraints distinct from adult-facing clinical LLMs; we are not aware of prior work fine-tuning age-stratified LoRA adapters for this setting (Section 2).

During development, an unexpected **configuration-ordering crossover** emerged (Standard better on 3B, Fast on 1B) – opposite existing rank-selection guidance at the 7B–70B scale – validated with a 110-run multi-seed campaign (Section 3.3) after an initial single-seed read proved non-reproducible.

The contributions of this paper are:

- *guard-bear*: a fine-tuned input-safety classifier showing its unmodified general-purpose base model performs *worse than random* on this population – resolved by fine-tuning to near-perfect, 5-seed-validated discrimination – evidence that safety classifiers cannot be assumed to transfer to a new population without dedicated evaluation.
- A complete age-adaptive pipeline, gated by an upstream safety classifier, trained end-to-end on consumer hardware; a prompted, non-fine-tuned baseline isolates fine-tuning’s real contribution as latency-compatible brevity, not readability (Section 3.1).
- An unexpected LoRA configuration-ordering crossover between Llama 3.2 1B and 3B, confirmed statistically real on both sides by a 110-run multi-seed campaign after an initial single-seed observation proved non-reproducible – though length-adjustment shows the 3B side is substantially a length effect, unlike 1B (Section 3.3).
- A rank sweep and layer sweep investigating the crossover’s mechanism (Section 3.4), and cross-architecture evidence (Qwen 2, SmolLM2, Qwen 2.5) that the Fast-vs-Standard directional pattern generalizes beyond Llama even where the rank-regime mechanism does not cleanly resolve.

## 2 Related Work

General-purpose prompt-injection classifiers such as **Prompt-Guard-86M** [1] target developer-facing override patterns; the closest child-safety effort, *KidRails* [6], fully fine-tunes an 8B response model without a separate classifier or baseline comparison. *guard-bear* instead evaluates a dedicated input classifier against its base model on a labeled pediatric threat taxonomy – to our knowledge, the first such evaluation for this population.

Prompted, not fine-tuned, LLMs have been evaluated as pediatric communication aids: Rao et al. [7] found off-the-shelf models can produce age-appropriate surgical explanations but called for further fine-tuning; a Norwegian evaluation similarly found generated child-directed language skewed too advanced [8]. *KidLM* [9] adapts masked-language-model pretraining to a children’s corpus but targets general modeling, not clinical generation. VR-based interventions reduce procedural anxiety [10] but are passive, scripted content; no prior work, to our knowledge, fine-tunes age-stratified LoRA adapters for pediatric hospital communication.

LoRA [4] and QLoRA [5] are standard for resource-constrained fine-tuning, but rank-selection guidance – including Biderman et al.’s [11] empirical study, the most thorough to date – is validated only at the 7B–70B scale; rank-ordering below 7B is uncharacterized by prior work. Our rank and layer sweeps (Section 3.4) address this gap.

## 3 Results

### 3.1 small-bear: Perplexity, Readability, Latency, and Style Separation

Table 1 reports a single seed-42 run of the four fine-tuned configurations and both base models; Section 3.3’s properly-powered campaign is the paper’s actual evidentiary basis for the crossover and is authoritative wherever the two differ. All four fine-tuned adapters substantially outperform zero-shot base models on  $FK \leq 7.0$  for the 5–11 group (Table 1: 56–66% vs. 12%) and inter-role style separation, not tabulated but computed identically to Table 5’s Style-Sep. column (Std-1B/3B, Fast-1B/3B: 0.950/0.930/0.960/0.920 vs. Base-1B/3B: 0.670/0.640) – though this classifier does not discriminate Standard from Fast, consistent with, not independent confirmation of, the readability pattern below. **But zero-shot is not the relevant alternative:** the same base model given one age-targeted system-prompt sentence<sup>1</sup>, no fine-tuning – single seed-42 run, one untuned prompt per role, not the multi-seed treatment applied to the crossover – reaches 92–98%  $FK \leq 7.0$  (Table 1, Prompted-1B/3B rows) and 0.970–1.000 style separation (same metric, not tabulated), matching or beating every adapter. It does not control length: 145–250 tokens vs. the adapters’ 88–124, pushing latency to 2.0–10.3s and leaving only 0–8% of turns under 1.0s vs. Fast-1B’s 62%. **Fine-tuning’s load-bearing contribution is decoupling register from length**, not readability alone, which prompting already gets for free. **Fast-1B is the only Llama configuration meeting the 1.0s VR latency target on average** (0.97s), though only 62% of its individual turns land under 1.0s (Table 1); all 3B configurations exceed 2.2s on average. Post-hoc, single-seed perplexity does not track FK-based readability cleanly – Fast is lower than Standard at *both* sizes (1B: 30.18 vs. 30.46; 3B: 22.91 vs. 24.88) – reported only as a construct-validity flag, not a confirmed finding absent multi-seed treatment.

On this seed, Fast-1B beats Standard-1B on FK pass rate (62% vs. 56%), while Standard-3B and Fast-3B tie (66% vs. 66%); Section 3.3’s campaign shows this single-run picture undersells the true, smaller 3B-side gap – the qualitative direction holds, but treat the margins here as illustrative, not as the paper’s statistical evidence.

### 3.2 guard-bear: Safety Classification Results

The clearest evidence that general-purpose safety tooling cannot be assumed to transfer to a new population comes from comparing *guard-bear* against unmodified Prompt-Guard-86M (threshold 0.50) on the 608-example held-out test set (Table 2), verified zero text overlap with training (Section 4.3). **The base model is worse than random:** ROC-AUC 0.448 is below chance, since its imperative-override feature space does not describe short, first-person pediatric queries or gibberish/out-of-scope/adult content. Its apparent 0.9385 recall is a saturation artifact, flagging 292 of 299 safe examples as unsafe too; the score is  $1 - P(\text{benign})$  from its native softmax, ruling out a label-flip artifact. **Fine-tuning resolves the discrimination failure consistently:** a 5-seed campaign (Table 2, rightmost column) gives ROC-AUC  $0.9992 \pm 0.0005$ ,

---

<sup>1</sup>Age 5–11: “You are Dr. Beary Good, a warm and friendly guide at a children’s hospital. You are talking to a child between the ages of 5 and 11. Use simple words, short sentences, and a warm, reassuring tone.” Age 12–18 substitutes “a teenager between the ages of 12 and 18. Use clear, respectful language appropriate for a teen audience.” Untuned, single prompt per role (Limitations).

**Table 1** Readability (50 age-matched held-out prompts/group) and latency/throughput (sequential, Mac Mini M4), single run (seed 42), current dataset vintage.  $FK \leq 7.0$  is calibrated to the younger band only. Prompted-{1B,3B} is the same zero-shot base model with a simple age-targeted system prompt (no fine-tuning) – the cheap alternative baseline discussed in Section 3.1.

| Config.     | Age   | FK   | SMOG | G.Fog | C-L  | TTR   | $FK \leq 7.0$ | Avg Lat.(s) | <1.0s |
|-------------|-------|------|------|-------|------|-------|---------------|-------------|-------|
| Standard-1B | 5–11  | 7.1  | 9.2  | 8.7   | 7.1  | 0.756 | 56%           | 1.12        | 26%   |
| Standard-1B | 12–18 | 10.3 | 12.7 | 13.0  | 10.7 | 0.734 | —             | 1.36        | 2%    |
| Standard-3B | 5–11  | 6.6  | 8.9  | 8.1   | 7.1  | 0.781 | 66%           | 2.22        | 0%    |
| Standard-3B | 12–18 | 10.4 | 12.8 | 12.9  | 10.6 | 0.733 | —             | 3.08        | 0%    |
| Fast-1B     | 5–11  | 6.5  | 9.0  | 8.0   | 6.8  | 0.752 | 62%           | 0.97        | 62%   |
| Fast-1B     | 12–18 | 9.7  | 12.2 | 12.1  | 10.0 | 0.725 | —             | 1.11        | 38%   |
| Fast-3B     | 5–11  | 6.5  | 9.0  | 8.1   | 7.1  | 0.724 | 66%           | 2.70        | 0%    |
| Fast-3B     | 12–18 | 10.0 | 12.5 | 12.5  | 10.4 | 0.663 | —             | 3.63        | 0%    |
| Base-1B     | 5–11  | 9.5  | 11.9 | 12.0  | 9.6  | 0.582 | 12%           | 1.90        | 16%   |
| Base-1B     | 12–18 | 10.4 | 12.9 | 13.1  | 11.6 | 0.553 | —             | 2.23        | 2%    |
| Base-3B     | 5–11  | 9.8  | 12.4 | 12.5  | 10.2 | 0.609 | 12%           | 4.54        | 8%    |
| Base-3B     | 12–18 | 11.7 | 13.6 | 14.7  | 11.4 | 0.551 | —             | 5.40        | 2%    |
| Prompted-1B | 5–11  | 4.5  | 7.1  | 5.8   | 5.2  | 0.612 | 98%           | 1.97        | 8%    |
| Prompted-1B | 12–18 | 7.4  | 10.0 | 9.3   | 7.7  | 0.569 | —             | 2.22        | 2%    |
| Prompted-3B | 5–11  | 4.7  | 7.5  | 6.2   | 5.6  | 0.661 | 92%           | 5.03        | 0%    |
| Prompted-3B | 12–18 | 8.4  | 11.3 | 10.6  | 8.8  | 0.581 | —             | 10.29       | 0%    |

every seed meeting all four deployment targets (recall/precision/F1/accuracy  $\geq 0.97/0.85/0.90/0.92$ ) – seed 42 is representative, not cherry-picked. A cheap TF-IDF+logistic-regression baseline, trained and tested on the identical split, is not trivial (0.9923 AUC – this task is not adversarially hard to separate lexically) but is measurably, consistently below the fine-tuned transformer on every metric, with roughly double guard-bear’s false negatives and false positives (Table 2): fine-tuning adds real discriminative value beyond bag-of-words features, not just clearing an easy bar. Largest subcategory gain over the base model: **gibberish\_malformed** (38%  $\rightarrow$  92%). A source-provenance check (safe all synthetic; unsafe 59% synthetic, 41% public corpora) finds 98.3% recall on synthetic-source unsafe vs. 100% pulled-source – too close to indicate provenance shortcut learning, though paraphrase leakage remains unchecked (Limitations). One caveat: the *threshold* varies widely by seed (0.94 to 0.04) though performance does not – always re-tune after retraining.

Table 1’s latency figures and *guard-bear*’s footprint together show a complete pipeline practical on one consumer machine; Fast-1B alone clears the 1.0s VR target, as do several cross-architecture configurations (Table 5). End-to-end latency with *guard-bear*’s pass included has not yet been measured (Limitations).

### 3.3 Statistical Validation of the Crossover: A Multi-Seed Campaign

The single-run pattern above (Section 3.1) was an unexpected finding uncovered during development, not a designed comparison, so it needed the same reproducibility scrutiny as any surprising result before being treated as real. The ablation and the rank/layer sweeps in Section 3.4 are single- or two-seed comparisons; Section 4.3 documents why that scrutiny was necessary – a 17-point swing on a nominally identical re-run, traced primarily to a dataset revision plus a smaller, unexplained non-determinism effect.

Given this, we resolved the crossover question directly with a dedicated campaign: retraining Standard and Fast at their exact defined configurations ( $r =$

**Table 2** guard-bear vs. base model and a simple lexical baseline, held-out test set (608 examples: 299 safe, 309 unsafe; deduplicated, verified zero text overlap with train/val). TF-IDF+LR: logistic regression on TF-IDF features, same train/test split. Rightmost column: 5-seed campaign (42, 123, 456, 789, 2024), mean  $\pm$  SD.

| Metric           | Base (t=0.50) | TF-IDF+LR | seed 42 | guard-bear (5-seed)                 |
|------------------|---------------|-----------|---------|-------------------------------------|
| Unsafe recall    | 0.9385        | 0.9612    | 0.9903  | <b>0.9883<math>\pm</math>0.0043</b> |
| Unsafe precision | 0.4983        | 0.9459    | 0.9903  | <b>0.9864<math>\pm</math>0.0062</b> |
| Unsafe F1        | 0.6510        | 0.9535    | 0.9903  | <b>0.9874<math>\pm</math>0.0049</b> |
| Accuracy         | 0.4885        | 0.9523    | 0.9901  | <b>0.9872<math>\pm</math>0.0050</b> |
| ROC-AUC          | 0.4477        | 0.9923    | 0.9998  | <b>0.9992<math>\pm</math>0.0005</b> |

| Confusion | Base | TF-IDF+LR | seed 42 | 5-seed sum ( $n=3040$ ) |
|-----------|------|-----------|---------|-------------------------|
| TN        | 7    | 282       | 296     | 1474                    |
| FP        | 292  | 17        | 3       | 21                      |
| FN        | 19   | 12        | 3       | 18                      |
| TP        | 290  | 297       | 306     | 1527                    |

8/num\_layers=16 and  $r = 4/\text{num\_layers}=8$ ) across many seeds, one consistent (current) dataset vintage throughout, role `age_5_11`,  $\text{FK} \leq 7.0$  pass rate –  $n = 25/\text{config}$  at 1B (pre-committed, no interim look); at 3B, an initial  $n = 15/\text{config}$  batch left the comparison marginal ( $t = 2.01$ ,  $p = 0.054$ ), so, per a plan fixed before unblinding stage two (a single top-up to  $n = 30$ , not an open-ended stopping rule), we collected 15 further seeds/config and combine both stages below via a pre-specified equal-weight inverse-normal combination test, the valid correction for this design.

**Table 3** Seed-campaign  $\text{FK} \leq 7.0$  pass rate (mean  $\pm$  SD across independent training runs).

| Config.     | $n$ | Mean  | SD  |
|-------------|-----|-------|-----|
| Fast-1B     | 25  | 63.4% | 6.3 |
| Standard-1B | 25  | 52.1% | 6.3 |
| Fast-3B     | 30  | 52.2% | 7.0 |
| Standard-3B | 30  | 58.3% | 6.9 |

On 1B, Fast beats Standard ( $t = 6.40$ ,  $p = 6.0 \times 10^{-8}$ ). On 3B, the naive pooled test gives  $t = 3.42$ ,  $p = 0.0011$ , but the interim look makes this anti-conservative; combining that stage with the independent 15-seed replication collected afterward ( $t = 2.75$ ,  $p = 0.010$ ) via the combination test gives the valid  $p = 0.0015$  – **the crossover survives the statistically correct test on both sides**. This campaign covers  $\text{FK} \leq 7.0$  pass rate only, for `age_5_11` only; the other four readability metrics, the `age_12_18` role, latency, and the classifier are single-run, as in Table 1, though now on the same consistent dataset vintage as this campaign.

Because Standard and Fast differ systematically in response length (Table 1), we tested whether the crossover survives controlling for length, across all 5,500 individual campaign responses. Fast is shorter than Standard at 1B (72 vs. 78 words,  $p = 4 \times 10^{-16}$ ) but longer at 3B (97 vs. 76 words,  $p = 9 \times 10^{-50}$ ). Residualizing FK grade on a

single word-count regression pooled across configs (an equal-slopes assumption we did not test) and comparing residuals by config, **the 1B crossover survives length-adjustment** (residual gap 0.35 FK grade levels,  $t = -5.95$ ,  $p = 3 \times 10^{-9}$ ), but **the 3B crossover is substantially a length effect** – the residual gap shrinks to 0.10 FK grade levels, only marginal ( $t = -1.92$ ,  $p = 0.055$ ). A joint logistic regression of pass/fail on config and word count agrees: the config coefficient stays large at 1B (0.35) but is near zero at 3B (0.007) once length is included. Standard’s 3B advantage is better described as producing shorter, not more simply-worded, responses.

### 3.4 Mechanism Sweeps: Rank and Layer Depth

The Standard/Fast configurations differ in both rank and `num.layers`, so even with the crossover confirmed (Section 3.3), it is not attributable to either factor alone. Our working hypothesis is capacity regularization (higher rank permits more overfitting on this LIMA-scale dataset) [4, 11], though prior work does not characterize sub-7B behavior. The sweeps use two seeds – lower power than the campaign above;  $\pm$ std is a two-point range, not a variance estimate. Both were retrained on the current vintage.

Fixing `num.layers=8` and sweeping  $r \in \{2, 4, 8, 16\}$  across four architectures (Llama 1B/3B, Qwen 2 0.5B, SmolLM2 360M; two seeds; 32 runs; Table 4, left), **only SmolLM2 360M shows a clean, monotonic rank-regime pattern** ( $48\% \rightarrow 63\%$ ), consistent with depth driving a preference for higher rank; Llama – the architecture the crossover claim is actually about – and Qwen are flat or only weakly increasing within noise, so this sweep does not mechanistically explain the paper’s own headline crossover for its primary model family. We treat SmolLM2’s sensitivity as the sweep’s one qualitatively consistent finding at  $n = 2/\text{rank}$ , and “rank is the operative variable” as suggestive, not resolved.

A companion layer sweep (fixed `rank=4`, `num.layers`  $\in \{4, 8, 16\}$ , Llama 1B/3B, two seeds, 12 runs; Table 4, right) replicates more consistently: **3B improves with depth** ( $57\% \rightarrow 63\%$ ) while **1B peaks at the fewest layers** (69% at layers=4, declining to 60% at layers=16), matching the original sweep’s qualitative conclusion. Latency scales with layer count for both sizes.

### 3.5 Cross-Architecture and Capacity-Ladder Validation

To test generalization beyond Llama, we evaluate Qwen 2 0.5B and SmolLM2 360M [12] and the Qwen 2.5 family (0.5B, 1.5B, 3B) for age 5–11 (Table 5), single-seed and exploratory. **SmolLM2 Standard dominates** (80% FK); **Qwen 2.5 shows Fast at or above Standard at every size**, suggesting no crossover for that family. The Style-Sep. column reflects the original, pre-revision vintage.

## 4 Methods

### 4.1 System Architecture

Raw text first passes through *guard-bear*, a binary input classifier; unsafe input is blocked before reaching *small-bear*. Safe input is routed by an age-gate, supplied by the calling VR application, to one of two LoRA adapters on a shared frozen 4-bit quantized base model, never fused into it. Emotional-support queries are in scope only as reassurance/redirection; clinical assessment and crisis intervention are out of scope.

**Table 4** Rank sweep (left; num\_layers=8 fixed) and layer sweep (right; rank=4 fixed) FK  $\leq 7.0$  pass rate, mean  $\pm$  std across two seeds (n=2; read as a range between the two runs, not a variance estimate), current dataset vintage.


| Model     | $r=2$       | $r=4$                       | $r=8$                      | $r=16$                      |
|-----------|-------------|-----------------------------|----------------------------|-----------------------------|
| Llama 1B  | 63 $\pm$ 10 | 59 $\pm$ 7                  | <b>64<math>\pm</math>6</b> | 62 $\pm$ 6                  |
| Llama 3B  | 54 $\pm$ 11 | 58 $\pm$ 3                  | 57 $\pm$ 1                 | <b>62<math>\pm</math>11</b> |
| Qwen 0.5B | 70 $\pm$ 3  | <b>71<math>\pm</math>13</b> | 71 $\pm$ 4                 | 60 $\pm$ 11                 |
| SmolLM2   | 48 $\pm$ 3  | 56 $\pm$ 6                  | 62 $\pm$ 6                 | <b>63<math>\pm</math>4</b>  |

| Model, layers | FK $\leq 7$ | Lat.(s)     | Tok/s        |
|---------------|-------------|-------------|--------------|
| 1B, 4         | <b>69%</b>  | 0.85        | <b>105.6</b> |
| 1B, 8         | 61%         | 0.84        | 97.4         |
| 1B, 16        | 60%         | 1.06        | 84.0         |
| 3B, 4         | 57%         | 2.73        | <b>45.0</b>  |
| 3B, 8         | 60%         | 2.67        | 43.2         |
| 3B, 16        | <b>63%</b>  | <b>2.24</b> | 39.5         |

**Table 5** Cross-architecture and Qwen 2.5 capacity-ladder results, age 5–11 (50 examples, seed 42, current dataset vintage; Style-Sep. is the inter-role style-separation classifier, 5-fold CV, chance = 0.50, †original pre-revision vintage, not yet re-validated).

| Config.           | FK $\leq 7.0$ | Avg Lat. | <1.0s | Tok/s or Tok | Style-Sep. <sup>†</sup>             |
|-------------------|---------------|----------|-------|--------------|-------------------------------------|
| Qwen 2 Std.       | 68%           | 0.54 s   | 100%  | 83 tok       | 0.940 $\pm$ 0.049                   |
| Qwen 2 Fast       | 62%           | 0.51 s   | 96%   | 94 tok       | 0.960 $\pm$ 0.058                   |
| SmolLM2 Std.      | 80%           | 0.87 s   | 72%   | 82 tok       | 0.950 $\pm$ 0.032                   |
| SmolLM2 Fast      | 60%           | 1.30 s   | 56%   | 137 tok      | 0.920 $\pm$ 0.040                   |
| Qwen2.5-0.5B Fast | 56%           | 0.45 s   | 100%  | 176.4        | 0.950 $\pm$ 0.032                   |
| Qwen2.5-0.5B Std. | 56%           | 0.57 s   | 100%  | 152.1        | 0.940 $\pm$ 0.097                   |
| Qwen2.5-1.5B Fast | 74%           | 1.07 s   | 58%   | 76.1         | <b>0.980 <math>\pm</math> 0.024</b> |
| Qwen2.5-1.5B Std. | 66%           | 1.28 s   | 6%    | 67.8         | 0.890 $\pm$ 0.073                   |
| Qwen2.5-3B Fast   | 70%           | 1.73 s   | 4%    | 43.0         | 0.970 $\pm$ 0.040                   |
| Qwen2.5-3B Std.   | 50%           | 2.18 s   | 0%    | 39.9         | 0.930 $\pm$ 0.068                   |

No system prompt is used, so the adapter weights encode register directly. Figure 1 shows the pipeline; the age\_5-11/Fast/1B branch meets the 1.0s VR latency target (Section 5).



**Fig. 1** System architecture. *guard-bear* (86M params, <50ms) screens input before an age-gate routes safe queries to one of two LoRA adapters on a shared frozen 4-bit Llama 3.2 base; adapters are never fused. Rank/layer and base size are runtime choices behind each adapter box, not drawn separately – Table 1’s four configurations are all instances of this diagram.

*guard-bear* classifies text as **safe** (a legitimate pediatric query, including benign out-of-scope curiosity) or **unsafe** (injection, jailbreak, adult/parental, gibberish, inappropriate, social engineering), a full parameter fine-tune of a small (86M) classifier, unlike *small-bear*. Design target: high recall, since a missed unsafe input is worse than a blocked legitimate one ( $\geq 0.97$ ).

## 4.2 Datasets

*small-bear*’s dataset was generated by Claude Opus 4.6 in extended thinking mode from the Children’s Hospital Patient and Family Guide [13], across five categories, each age-stratified (5–11, 12–18). An initial 1000-example release was superseded by a quality-curation pass adding a training-only **edge\_cases** category (50) and 73 diversified examples, for 1123 total, split 562/561, and 100 new validation examples; Section 4.3 documents how this revision affected earlier results. Examples are instruction/response pairs, without a system prompt, so the model encodes register in its weights.

*guard-bear*’s dataset combines synthetic safe/unsafe examples with unsafe examples from public prompt-injection and jailbreak datasets: 4,089 raw examples, deduplicated to 4,053 unique (2,055 unsafe, 1,998 safe) across 14 subcategories (taxonomy in the repository). “Safe” means *not actively harmful or adversarial*, not “in *small-bear*’s scope”; out-of-scope questions are handled downstream. Dedup precedes splitting, so the 2,836/609/608 train/val/test split (Section 4.3) is disjoint by exact text.

## 4.3 Base Models, LoRA Configuration, and Training Protocol

We use the Llama 3.2 Instruct family [3] at 1B/3B, 4-bit quantized via `mlx_lm` [14] on Apple Silicon; full-parameter fine-tuning of the 3B model exceeds the 16GB ceiling, making LoRA the only viable strategy. *guard-bear* fine-tunes all 86M parameters via the Metal Performance Shaders backend.

LoRA [4] freezes a pre-trained weight matrix  $W_0 \in \mathbb{R}^{d \times k}$  and expresses updates as  $\Delta W = BA$ , applied to all linear projections in the targeted layers with `scale=20.0` held constant (a raw output multiplier, not an alpha/rank ratio). **Standard** ( $r = 8$ , `num_layers=16`) covers the upper 16 layers; **Fast** ( $r = 4$ , `num_layers=8`) covers the upper 8. The two differ in both rank and layer count simultaneously; Sections 3.3–3.4 confirm the resulting crossover is real.

*small-bear*’s adapters train with Adam, cosine-decay LR (peak  $1 \times 10^{-5}$ , 100-step warmup, 600 steps), batch size 4, max sequence length 768, dropout 0.0, loss masked to assistant tokens; all eight runs use seed 42. *guard-bear* fine-tunes **meta-llama/Prompt-Guard-86M** [1] (DeBERTa-v2-based [15]) on the full training split, threshold tuned post-hoc by sweeping [0.01, 0.95] and selecting the best-F1 point meeting  $\geq 0.97$  recall on the validation split, giving **0.94**.

**Reproducibility and dataset vintage.** A re-run of the nominal Fast-1B configuration (identical model, data, seed) produced an  $\text{FK} \leq 7.0$  pass rate 17 points lower than an earlier draft’s single run. A library-version confound was tested and found insufficient on its own; the training/validation data had also been revised between runs (Section 4.2), and re-scoring the original adapter on the revised validation set alone closed most of the gap. That test is itself imperfect – run against the current, not period-matched, validation set – so it was never cleanly isolated from the vintage

confound either. A separate rerun (old-vintage data held fixed throughout, only `mlx` version differing) showed its own genuine, unexplained 16–18-point swing – distinct from the Fast-1B gap above, evidence of real run-to-run instability independent of dataset vintage. The rank sweep, layer sweep, and initial cross-architecture results were found to predate the revision too and were retrained on the current vintage (Table 5’s Style-Sep. column excepted, still provisional), motivating the dedicated multi-seed campaign in Section 3.3: a single-run result cannot be trusted without independent scrutiny.

*guard-bear* received the same scrutiny: splits were audited for exact-text overlap, the dataset deduplicated before splitting to guarantee disjointness by construction, and the model validated with a 5-seed campaign, not a single run (Section 3.2). One wrinkle is worth recording precisely: an initial overlap check compared validation against the full train+val pool it was drawn from, not the actual disjoint training set, making leakage look far more severe (612 of 614) than the true figure ( $\sim 36$  exact-duplicate rows among 4,089 raw examples caused 10 of 614 validation and 3 of 613 test examples to genuinely overlap) – both the measurement error and the real, smaller leak are now fixed at the root.

#### 4.4 Evaluation Framework

For *small-bear*: **readability** via `textstat` [16] (Flesch-Kincaid [17] primary, pass/fail at  $FK \leq 7.0$ ; SMOG [18], Gunning Fog [19], Coleman-Liau [20], type-token ratio); **latency**; **style separation** via TF-IDF+logistic-regression (scikit-learn [21], 5-fold CV, chance = 0.50).  $FK \leq 7.0$  is a general patient-material health-literacy bar, not a per-age-band fit – looser than the  $\sim$ grade 5–6 level of the youngest 5–11 patients – used as one shared cutoff across configurations, not a claim of exact age-fit. These are proxies, not direct measures, of age-adapted communication (Section 5). For *guard-bear*: unsafe recall, precision, F1, accuracy, ROC-AUC, and per-subcategory recall.

**Ethical approval declarations.** Not applicable in the formal IRB/human-subjects sense: no experiments on live vertebrates, humans, or human tissue; all data are synthetic or public reference/prompt-injection materials, with no patient contact or identifying information. This does not certify the deployment context: both systems target a real pediatric population, yet no clinician has reviewed any output (Limitations).

### 5 Discussion

For *guard-bear*, the central finding is that **a safety classifier cannot be assumed to transfer to a new deployment population without evaluation**; the gap is severe, worse-than-random, not suboptimal thresholding. *small-bear* without an input safety gate is not, by itself, deployable; each component’s latency alone is practical on one consumer machine (Section 3.2), though not yet measured combined.

**Fast-1B is the best Llama-family configuration** for VR deployment (Figure 1): the only option meeting the 1.0s latency target on average with strong readability and style separation; SmolLM2 Standard (Section 3.5) beats it on FK pass rate and is a strong alternative where architecture is unconstrained. The **crossover** (Section 3.3) suggests small-model LoRA configuration below 7B does not follow

rank-selection guidance developed at larger scales, though only the 1B side is length-independent and its mechanism is only partially resolved (Section 3.4) – necessary but not sufficient for genuine age-adapted communication, since a response can be short and lexically distinct while still emotionally mismatched, a gap only human evaluation can close.

## 6 Conclusion

This paper shows two things: a general-purpose input-safety classifier does not transfer to a new, vulnerable population without dedicated evaluation (**Prompt-Guard-86M**’s worse-than-random 0.448 AUC, resolved to  $0.9992 \pm 0.0005$  across a 5-seed campaign), and a complete age-adaptive pipeline is achievable on one consumer machine – though a prompted baseline reaches the same readability, fine-tuning’s real edge is reaching it without the length cost that breaks real-time latency (Abstract).

We also uncovered an unexpected LoRA configuration-ordering crossover between 1B and 3B Llama models (Section 3.3), surviving an initial non-reproducible single-seed observation – though only the 1B side is a length-independent register effect, the 3B side being substantially a length effect ( $p = 0.055$  once controlled); a rank/layer sweep attributes the crossover partly to rank and depth, though the rank mechanism replicates cleanly only for SmolLM2 – existence established, precise cause only partially so. Neither component alone is deployable: *small-bear* presumes benign input, *guard-bear* makes that presumption safe.

Limitations: Section 4.3 documents this pipeline’s non-determinism and dataset-vintage confound, accounted for above; one piece remains provisional (Table 5’s Style-Sep. column). The rank sweep’s mechanism claim remains qualified at two-seed power (Section 3.4, Conclusion). Length-adjustment was untested on the sweeps’ architectures; single-seed perplexity does not track the crossover; models above 3B remain untested. Training and validation data share one generative process, and every metric is fully automated: no human rater has judged any output, so this closed loop confirms internal, not external, validity, and neither component has clinical validation. Reported latency excludes *guard-bear*’s pass; its dataset has not been red-teamed, and paraphrase-level (as opposed to exact-text) leakage remains unchecked despite the multi-seed campaign (Section 3.2). **Most operationally consequential:** that campaign’s per-seed threshold ranges from 0.94 to 0.04 though performance does not – a fixed threshold cannot be assumed to transfer to a retrained model, and deployment must re-tune on fresh validation data after every retrain, which may not always be available. A blocked query returns no response at all (Figure 1); for a distressed child, silence instead of a graceful redirect is itself an unaddressed harm. The prompted-baseline comparison (Section 3.1) uses one untuned system prompt per role; a different prompt could shift its readability/latency tradeoff. Bundling three contributions (safety classifier, response-model ablation, mechanism sweeps) under one page budget leaves each with less room than it would have standalone – most visibly, *guard-bear*’s related work is a single paragraph, and the sweeps run at two seeds rather than the response-model campaign’s 25–30.

**Supplementary information.** Validation-set outputs, dataset files, and scripts are available at the repositories under Data availability below; no separate bundle accompanies this submission.

**Acknowledgements.** Not applicable.

## Declarations

- **Funding:** No funding to disclose.
- **Conflict of interest/Competing interests:** The authors declare no conflicts of interest.
- **Ethics approval and consent to participate; consent for publication:** Not applicable (no human subjects, patient data, or clinical trial); all data are synthetic or public reference/prompt-injection materials. Not a certification of clinical or deployment readiness – no clinician has reviewed any output (Limitations).
- **Data, materials, and code availability:** *small-bear*’s code/data at <https://github.com/abarros6/small-bear>; *guard-bear*’s code/data at <https://github.com/abarros6/guard-bear>.
- **Author contribution:** S.d.R. and R.E. supervised the research. A.B. ran all experiments, designed *guard-bear*, and wrote the paper except the Introduction, written by K.K. All authors reviewed and approved the final manuscript.

## References

- [1] Meta AI: Prompt-Guard-86M. <https://huggingface.co/meta-llama/Prompt-Guard-86M> (2024)
- [2] Zhou, C., Liu, P., Xu, P., Iyer, S., Sun, J., Maddela, R., Peng, X., Shrivastava, S., Lewis, M., Zettlemoyer, L., Levy, O.: LIMA: Less is more for alignment. In: Advances in Neural Information Processing Systems (NeurIPS) (2023)
- [3] Dubey, A., et al.: The llama 3 herd of models. arXiv preprint arXiv:2407.21783 (2024). Meta AI
- [4] Hu, E.J., Shen, Y., Wallis, P., Allen-Zhu, Z., Li, Y., Wang, S., Wang, L., Chen, W.: LoRA: Low-rank adaptation of large language models. In: Proc. Int. Conf. Learning Representations (ICLR) (2022)
- [5] Dettmers, T., Pagnoni, A., Holtzman, A., Zettlemoyer, L.: QLoRA: Efficient finetuning of quantized LLMs. In: Advances in Neural Information Processing Systems (NeurIPS) (2023)
- [6] Arcee AI: KidRails: Replicating a Child-Safe LLM Model. <https://github.com/arcee-ai/KidRails> (2025)
- [7] Rao, A.S., Mazumder, A., Roux, E., Young, C., Bott, E., Wang, J., Kochis, M., Stetson, A., Butler, A., Hilker, S., Succi, M.D.: Designing patient-centered communication aids in pediatric surgery using large language models. Journal of Pediatric Surgery **60**(12), 162654 (2025)
- [8] Hassan, S.Z., Halvorsen, P., Johnson, M.S., Lison, P.: Evaluating LLMs on generating age-appropriate child-like conversations. arXiv preprint arXiv:2510.24250

(2025)

- [9] Nayeem, M.T., Rafiei, D.: KidLM: Advancing language models for children – early insights and future directions. In: Proc. 2024 Conf. Empirical Methods in Natural Language Processing (EMNLP), Miami, FL, USA, pp. 4813–4836 (2024)
- [10] Ryu, J.-H., Park, J.-W., Choi, S.I., Kim, J.Y., Lee, H., Yoo, H.-J., Han, S.-H.: Virtual reality vs. tablet video as an experiential education platform for pediatric patients undergoing chest radiography: A randomized clinical trial. *J. Clin. Med.* **10**(11), 2486 (2021)
- [11] Biderman, D., Ortiz, J.G., Portes, J., Paul, M., Greengard, P., Jennings, C., King, D., Havens, S., Chiley, V., Frankle, J., Blakeney, C., Cunningham, J.P.: LoRA learns less and forgets less. *Transactions on Machine Learning Research (TMLR)* (2024)
- [12] Allal, L.B., et al.: SmolLM2: When smol goes big. arXiv preprint arXiv:2502.05737 (2025). Hugging Face
- [13] Children’s Hospital London Health Sciences Centre: Children’s Hospital Patient and Family Guide. <https://www.lhsc.on.ca/childrens-hospital/childrens-hospital-patient-and-family-guide>. Accessed Mar. 26, 2026
- [14] Apple: MLX: Efficient and Flexible Machine Learning on Apple Silicon. <https://github.com/ml-explore/mlx> (2023)
- [15] He, P., Liu, X., Gao, J., Chen, W.: DeBERTa: Decoding-enhanced BERT with disentangled attention. In: Proc. Int. Conf. Learning Representations (ICLR) (2021)
- [16] Bansal, S.: textstat: Python Library to Calculate Readability Statistics of a Text Object. <https://pypi.org/project/textstat/>. v0.7.13 (2023)
- [17] Kincaid, J.P., Fishburne, R.P., Rogers, R.L., Chissom, B.S.: Derivation of new readability formulas for navy enlisted personnel. Technical Report Research Branch Report 8-75, Naval Technical Training Command (1975)
- [18] McLaughlin, G.H.: SMOG grading — a new readability formula. *J. Reading* **12**(8), 639–646 (1969)
- [19] Gunning, R.: *The Technique of Clear Writing*. McGraw-Hill, New York, NY, USA (1952)
- [20] Coleman, M., Liau, T.L.: A computer readability formula designed for machine scoring. *J. Applied Psychology* **60**(2), 283–284 (1975)
- [21] Pedregosa, F., et al.: Scikit-learn: Machine learning in python. *J. Machine Learning Research* **12**, 2825–2830 (2011)